



UNIVERSITÀ
CATTOLICA
del Sacro Cuore

Kubernetes Details



INTRO TO KUBERNETES



kubernetes



KEY CONCEPTS

Container orchestrator

Runs and manages containers

Unified API for deploying web applications, batch jobs, and databases

Maintains and tracks the global view of the cluster

Supports multiple cloud and bare-metal environments

Manage applications, not machines

Specify the application to deploy, let Kubernetes decide what machines to run it on

Designed for extensibility

Implemented as resources, add custom resource types

Rich ecosystem of plug-ins for scheduling, storage, networking

Open source project managed by the Cloud Native Computing Foundation in the Linux Foundation

Inspired and informed by Google's experiences and internal systems

100% open source, written in Go



KEY CONCEPTS

KUBERNETES ARCHITECTURE

Master nodes

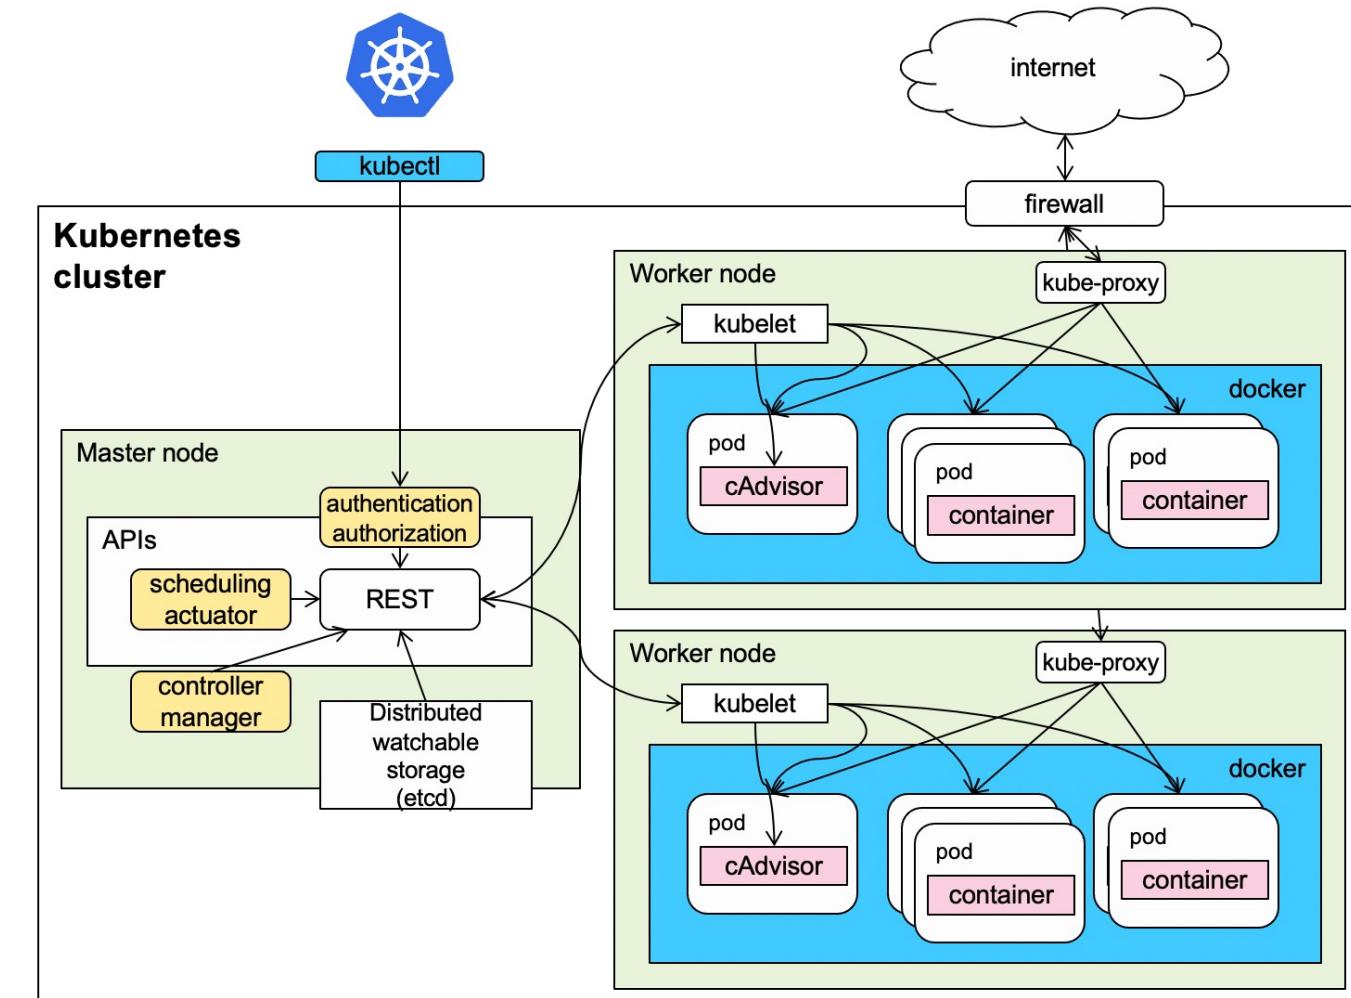
Computer with processes that manage the cluster

Multiple nodes for High Availability

Worker nodes

Computers that host pods

Pods host application containers





KEY CONCEPTS

Cluster

Collection of worker nodes managed by the same master node

Makes the worker nodes behave like one big computer

Master node

Controls and manages the cluster

Scheduling and replication logic

Generally two or more master nodes for resiliency, but are not used for scaling out the cluster

Worker node

Node where workloads run in pods

Docker – Docker engine for running containers

kubelet

Kubernetes agent that accepts commands from the master

Manages pods in the node

cAdvisor – Container Advisor provides resource usage and performance statistics

kube-proxy – network proxy service responsible for routing activities for inbound or ingress traffic

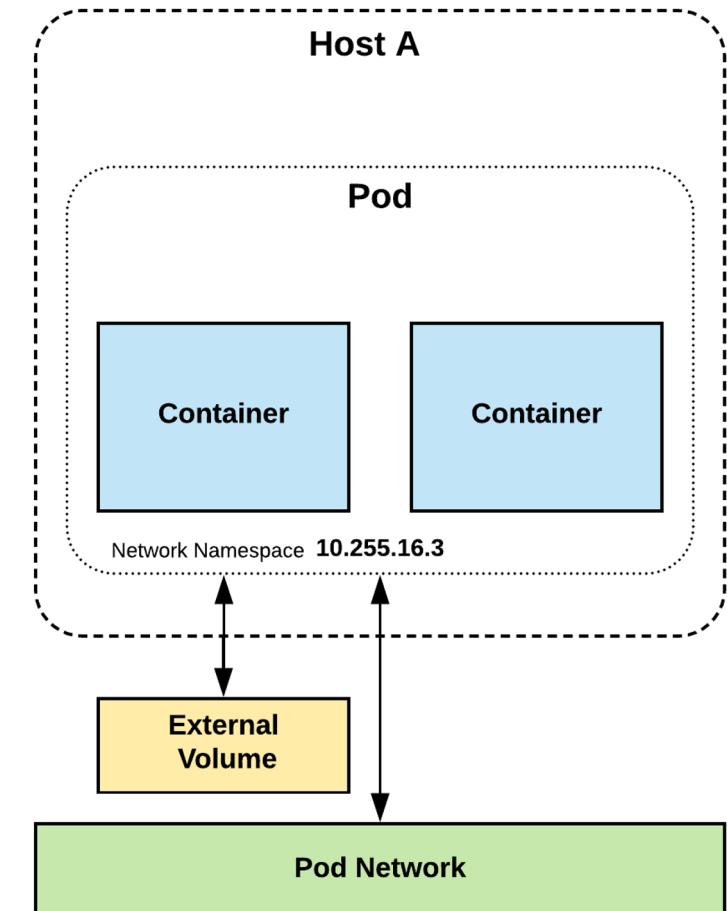


KEY CONCEPTS

POD

Atomic unit or smallest “*unit of work*” of Kubernetes.

Pods are **one or MORE containers** that share volumes, a network namespace, and are a part of a **single context**.





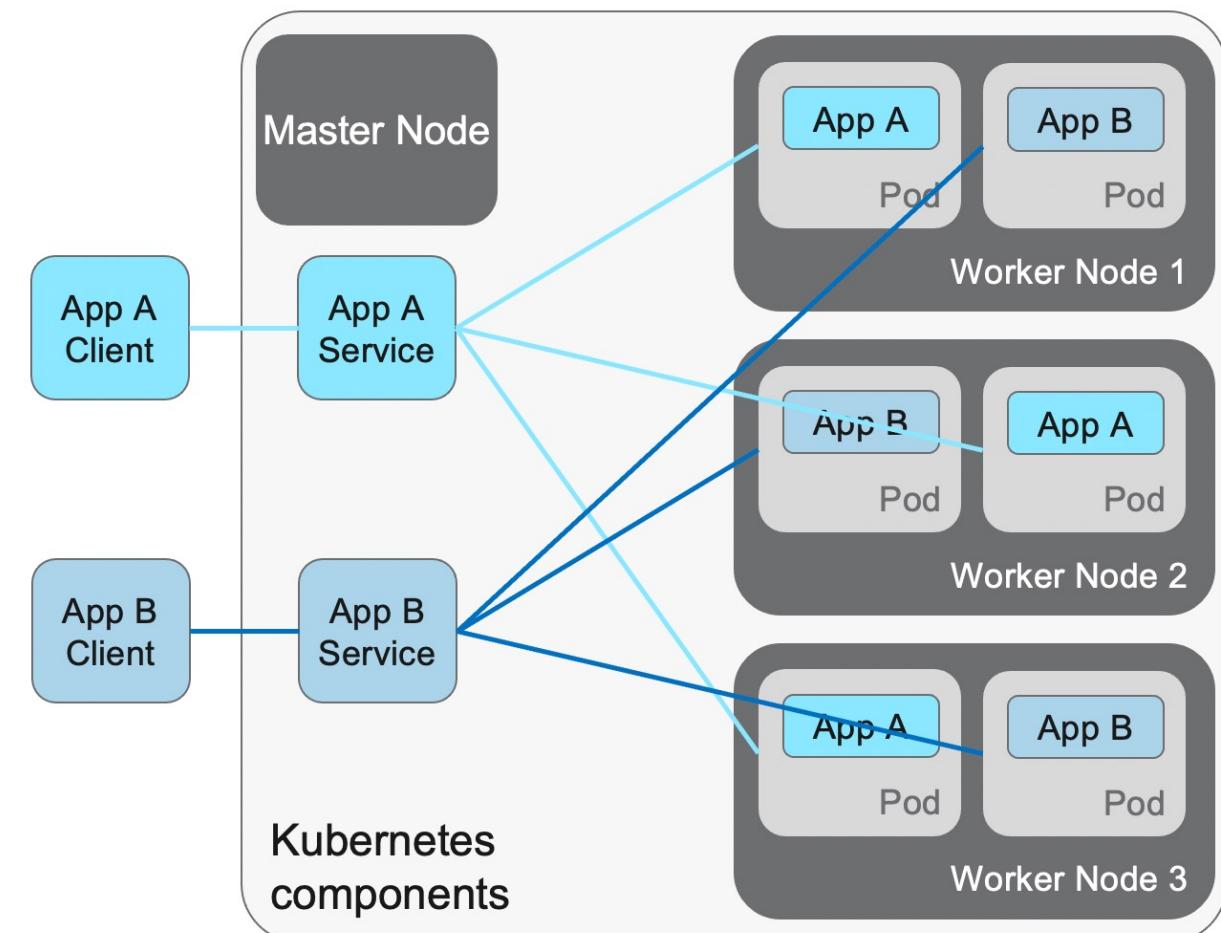
KEY CONCEPTS

KUBERNETES ARCHITECTURE: WORKLOADS

Container
Packaging of an app

Pod
Unit of deployment

Service
Fixed endpoint for 1+ pods





KEY CONCEPTS

Container

Unit of packaging

Pod

Smallest deployment unit in Kubernetes

A collection of containers that run on a worker node

Each pod has its own IP

A pod shares a PID namespace, network, and hostname

Service

Collection of pods exposed as an endpoint

Information stored in the Kubernetes cluster state and networking info propagated to all worker nodes

Types

ClusterIP – Exposes the service on a cluster-internal IP

NodePort – Exposes the service on each Node's IP at a static port

LoadBalancer – Exposes the service externally using a cloud provider's load balancer



KEY CONCEPTS

KUBERNETES TERMINOLOGY: DEPLOYMENT

Deployment

A set of pods to be deployed together, such as an application

Declarative: Revising a Deployment creates a ReplicaSet describing the desired state

Rollout: Deployment controller changes the actual state to the desired state at a controlled rate

Rollback: Each Deployment revision can be rolled back

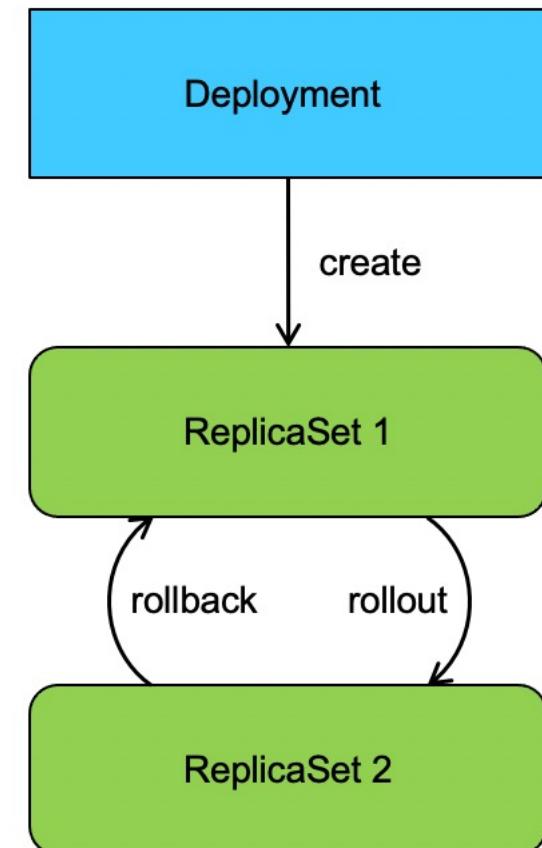
Scale and autoscale: A Deployment can be scaled

ReplicaSet

A set of pod templates that describe a set of pod replicas

Uses a template that describes specifically what each pod should contain

Ensures that a specified number of pod replicas are running at any given time





KEY CONCEPTS

KUBERNETES TERMINOLOGY: NAMING

Name

Each resource object by type has a unique name

Namespace

Resource isolation: Each namespace is a virtual cluster within the physical cluster

Resource objects are scoped within namespaces

Low-level resources are not in namespaces: nodes, persistent volumes, and namespaces themselves

Names of resources need to be unique within a namespace, but not across namespaces

Resource quotas: Namespaces can divide cluster resources

Initial namespaces

`default` – The default namespace for objects with no other namespace

`kube-system` – The namespace for objects created by the Kubernetes system

Resource Quota

Limits resource consumption per namespace

Limit can be number of resource objects by type (pods, services, etc.)

Limit can be total amount of compute resources (CPU, memory, etc.)

Overcommit is allowed; contention is handled on a first-come, first-served basis



KEY CONCEPTS

KUBERNETES CONFIGURING RESOURCES AND CONTAINERS

Label

Metadata assigned to Kubernetes resources (pods, services, etc.)

Key-value pairs for identification

Critical to Kubernetes as it relies on querying the cluster for resources that have certain labels

Selector

An expression that matches labels to identify related resources

ConfigMap

Configuration values to be used by containers in a pod

Stores configuration outside of the container image, making containers more reusable

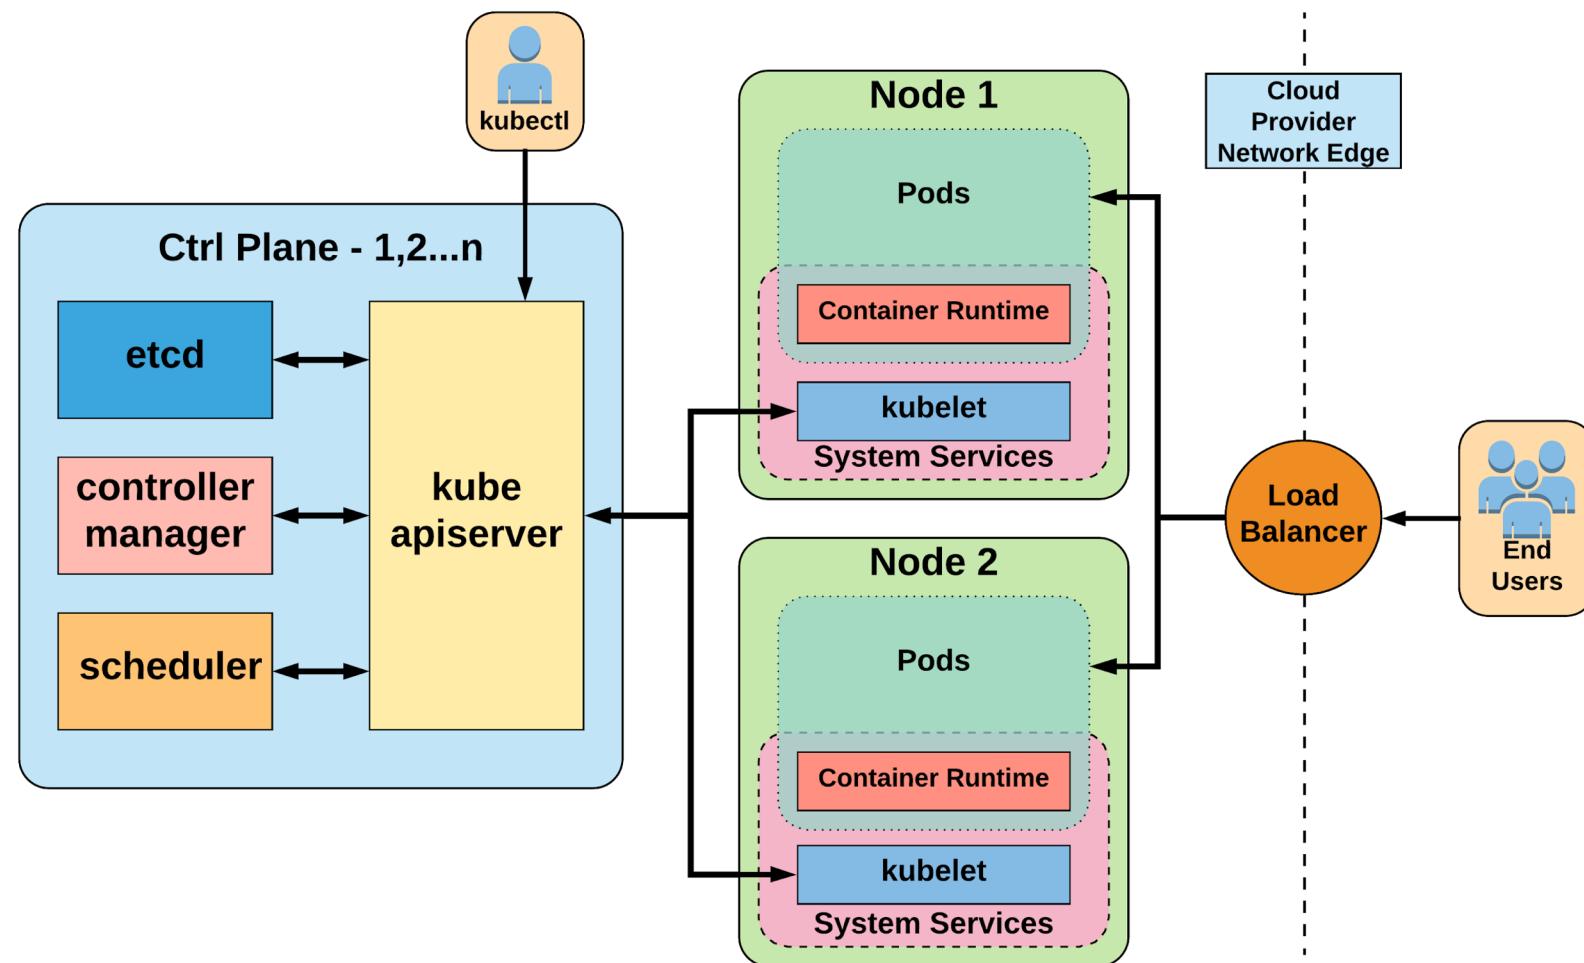
Secrets

Sensitive info that containers need to read or consume

Encrypted in special volumes mounted automatically



ARCHITECTURE RECAP





WORKING WITH KUBERNETES

OBJECT EXPRESSION - YAML

Files or other representations of Kubernetes Objects are generally represented in YAML.

A “*Human Friendly*” data serialization standard.

Uses white space (specifically spaces) alignment to denote ownership.

Three basic data types:

mappings - hash or dictionary,

sequences - array or list

scalars - string, number, boolean etc

..THERE IS ALSO A COMMAND LINE INTERFACE (CLI)



WORKING WITH KUBERNETES

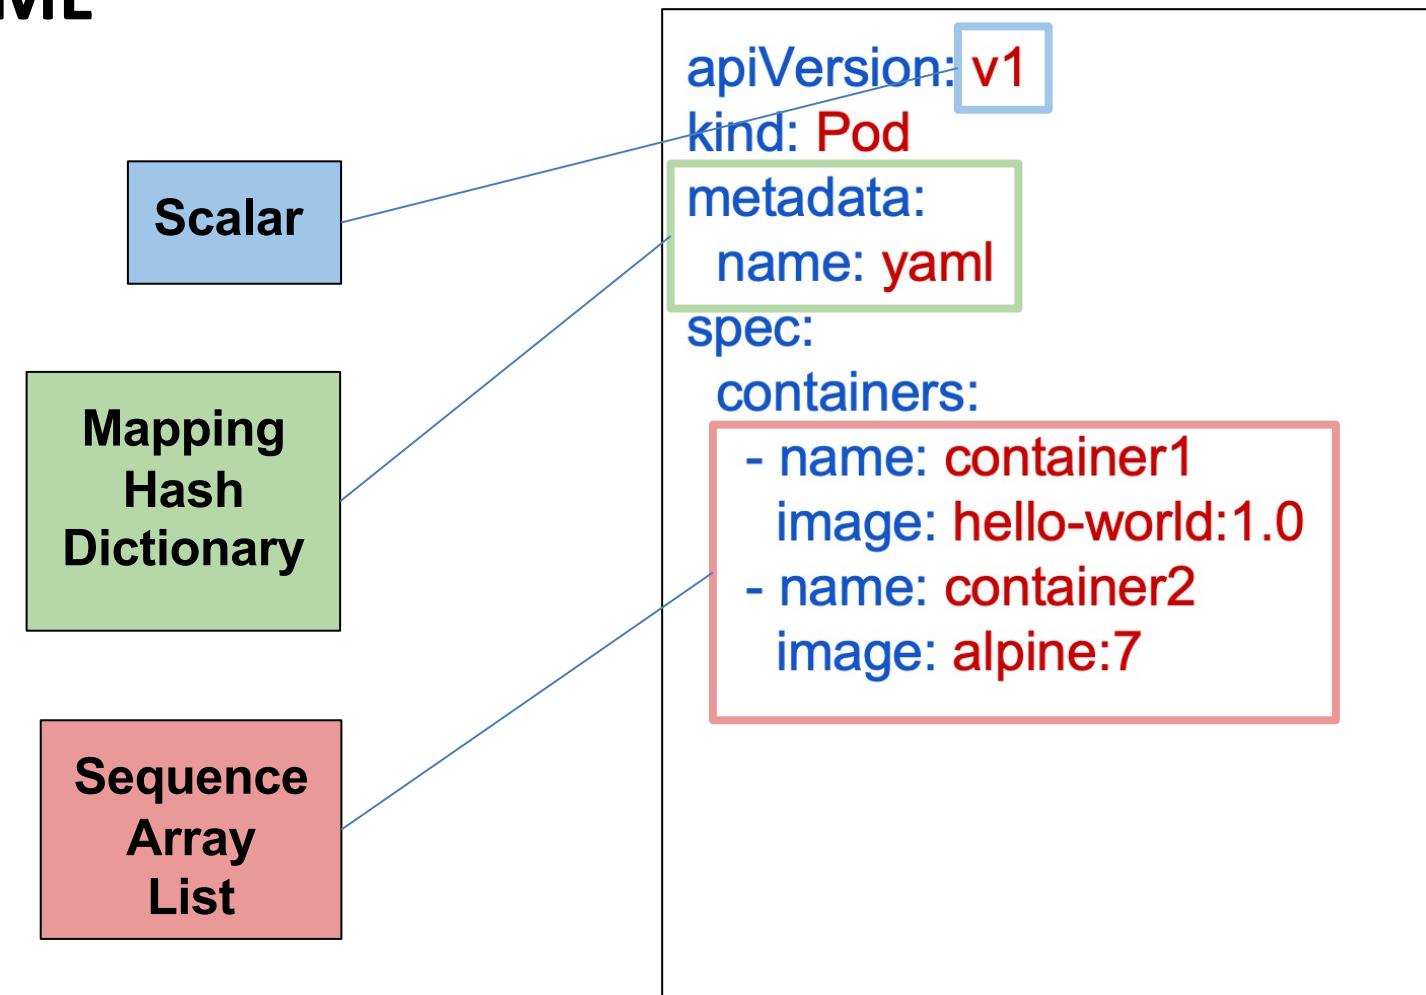
OBJECT EXPRESSION - YAML

```
apiVersion: v1
kind: Pod
metadata:
  name: yaml
spec:
  containers:
    - name: container1
      image: hello-world:1.0
    - name: container2
      image: alpine:7
```



WORKING WITH KUBERNETES

OBJECT EXPRESSION - YAML





WORKING WITH KUBERNETES

YAML

```
apiVersion: v1
kind: Pod
metadata:
  name: pod-example
spec:
  containers:
    - name: nginx
      image: nginx:stable-alpine
      ports:
        - containerPort: 80
```

JSON

```
{
  "apiVersion": "v1",
  "kind": "Pod",
  "metadata": {
    "name": "pod-example"
  },
  "spec": {
    "containers": [
      {
        "name": "nginx",
        "image": "nginx:stable-alpine",
        "ports": [ { "containerPort": 80 } ]
      }
    ]
  }
}
```



WORKING WITH KUBERNETES

POD CONTAINER ATTRIBUTES

- **name** - The name of the container
- **image** - The container image
- **ports** - array of ports to expose. Can be granted a friendly name and protocol may be specified
- **env** - array of environment variables
- **command** – Entrypoint array (equiv to Docker **ENTRYPOINT**)
- **args** - Arguments to pass to the command (equiv to Docker **CMD**)

Container

```
name: nginx
image: nginx:stable-alpine
ports:
  - containerPort: 80
    name: http
    protocol: TCP
env:
  - name: MYVAR
    value: isAwesome
command: ["/bin/sh", "-c"]
args: ["echo ${MYVAR}"]
```



WORKING WITH KUBERNETES

SERVICES

Unified method of accessing the exposed workloads of Pods.

Durable resource (unlike Pods)

static cluster-unique IP

static namespaced DNS name

<service name>. <namespace>. svc.cluster.local



WORKING WITH KUBERNETES

SERVICES

There are 3 major service types:

ClusterIP (default)

NodePort

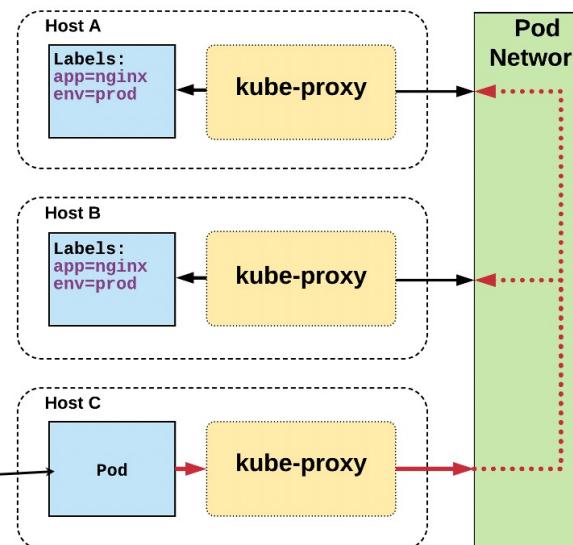
LoadBalancer



WORKING WITH KUBERNETES

SERVICES

ClusterIP services exposes a service on a strictly cluster internal virtual IP.



```
/ # nslookup example-prod.default.svc.cluster.local  
Name:   example-prod.default.svc.cluster.local  
Address 1: 10.96.28.176 example-prod.default.svc.cluster.local
```

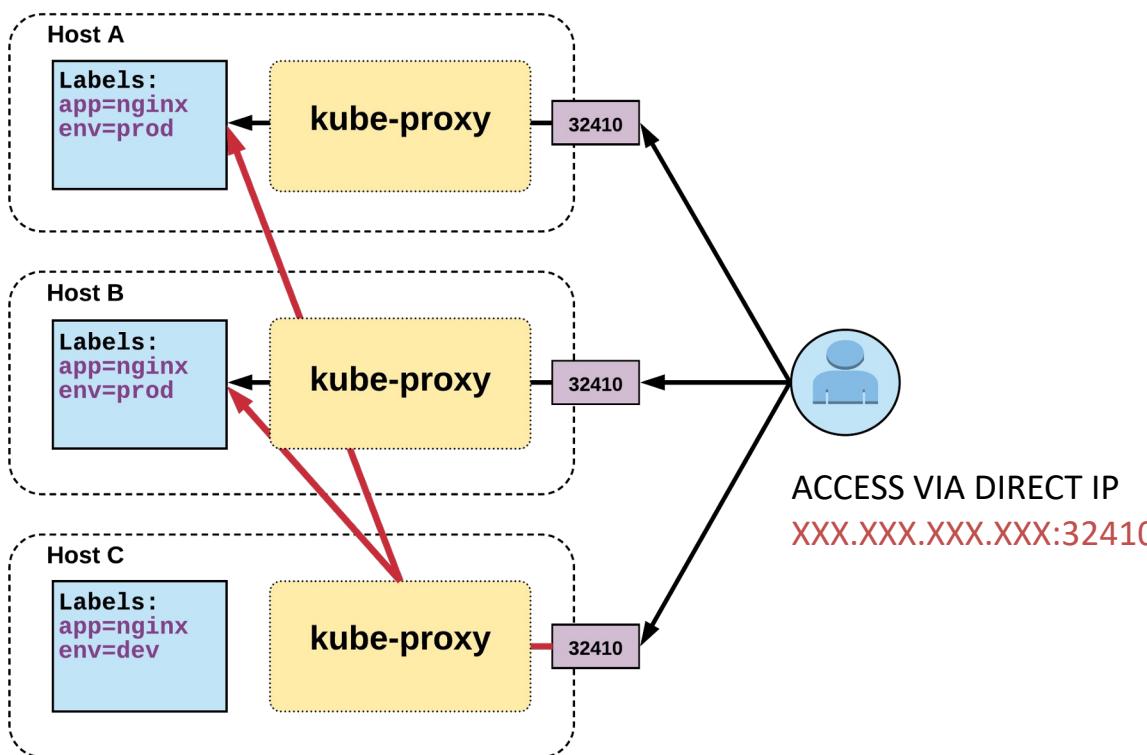
```
apiVersion: v1  
kind: Service  
metadata:  
  name: example-prod  
spec:  
  selector:  
    app: nginx  
    env: prod  
  ports:  
    - protocol: TCP  
      port: 80  
      targetPort: 80
```



WORKING WITH KUBERNETES

SERVICES

NodePort Exposes a port on every node's IP.



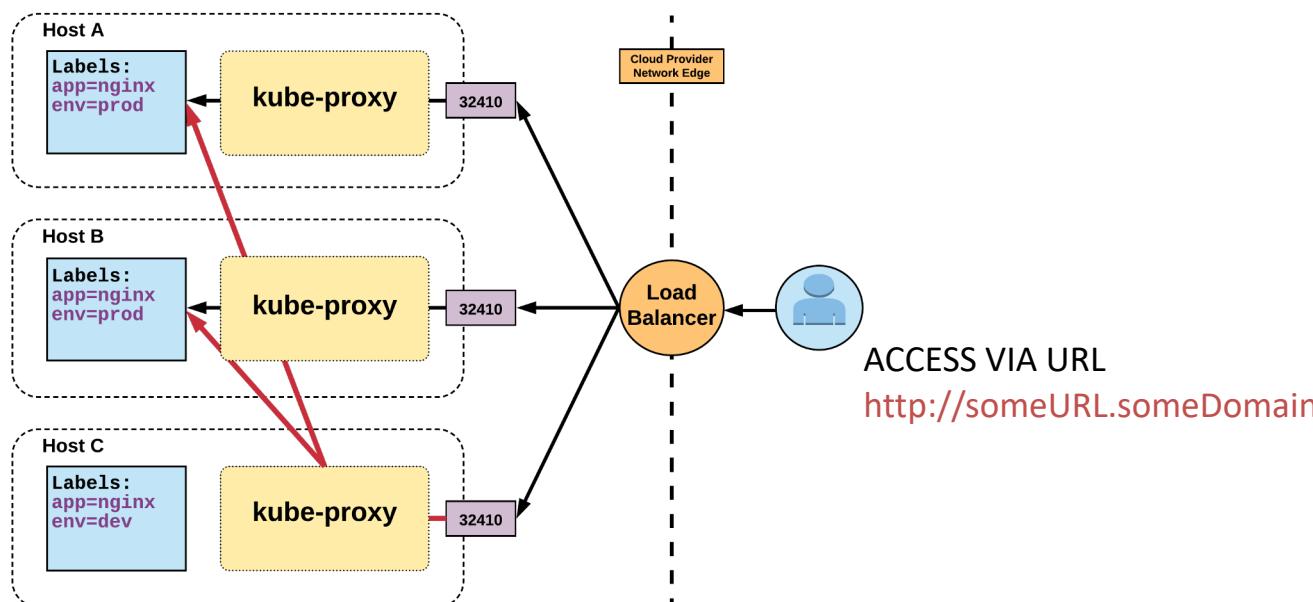
```
apiVersion: v1
kind: Service
metadata:
  name: example-prod
spec:
  type: NodePort
  selector:
    app: nginx
    env: prod
  ports:
    - nodePort: 32410
      protocol: TCP
      port: 80
      targetPort: 80
```



WORKING WITH KUBERNETES

SERVICES

LoadBalancer Works in conjunction with an external system to map a cluster external IP to the exposed service.



```
apiVersion: v1
kind: Service
metadata:
  name: example-prod
spec:
  type: LoadBalancer
  selector:
    app: nginx
    env: prod
  ports:
    protocol: TCP
    port: 80
    targetPort: 80
```



WORKING WITH KUBERNETES

SERVICES

ExternalName is used to reference endpoints **OUTSIDE** the cluster.

```
apiVersion: v1
kind: Service
metadata:
  name: example-prod
spec:
  type: ExternalName
spec:
  externalName: example.com
```



DEPLOYMENT

- **revisionHistoryLimit**: The number of previous iterations of the Deployment to retain.
- **strategy**: Describes the method of updating the Pods based on the **type**. Valid options are **Recreate** or **RollingUpdate**.
 - **Recreate**: All existing Pods are killed before the new ones are created.
 - **RollingUpdate**: Cycles through updating the Pods according to the parameters: **maxSurge** and **maxUnavailable**.

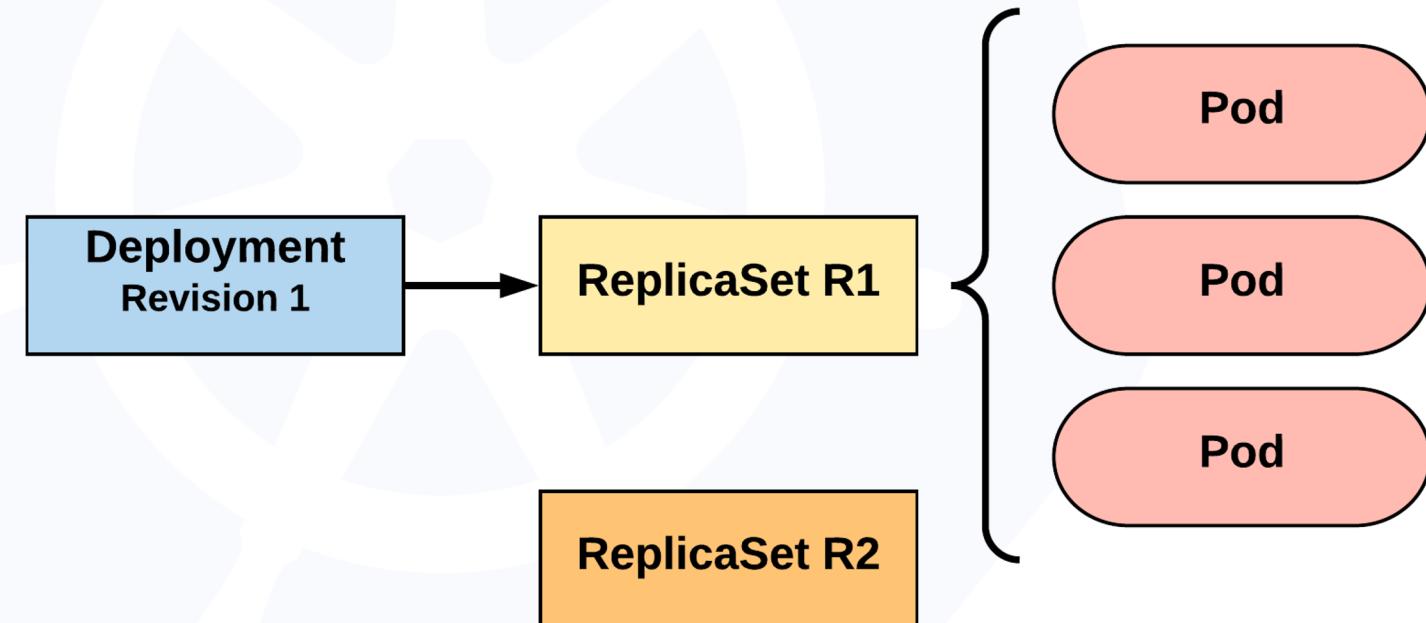
```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: deploy-example
spec:
  replicas: 3
  revisionHistoryLimit: 3
  selector:
    matchLabels:
      app: nginx
      env: prod
  strategy:
    type: RollingUpdate
    rollingUpdate:
      maxSurge: 1
      maxUnavailable: 0
  template:
    <pod template>
```



ROLLING UPDATE CONCEPT

Updating pod template generates a new **ReplicaSet** revision.

R1 pod-template-hash:
676677fff
R2 pod-template-hash:
54f7ff7d6d

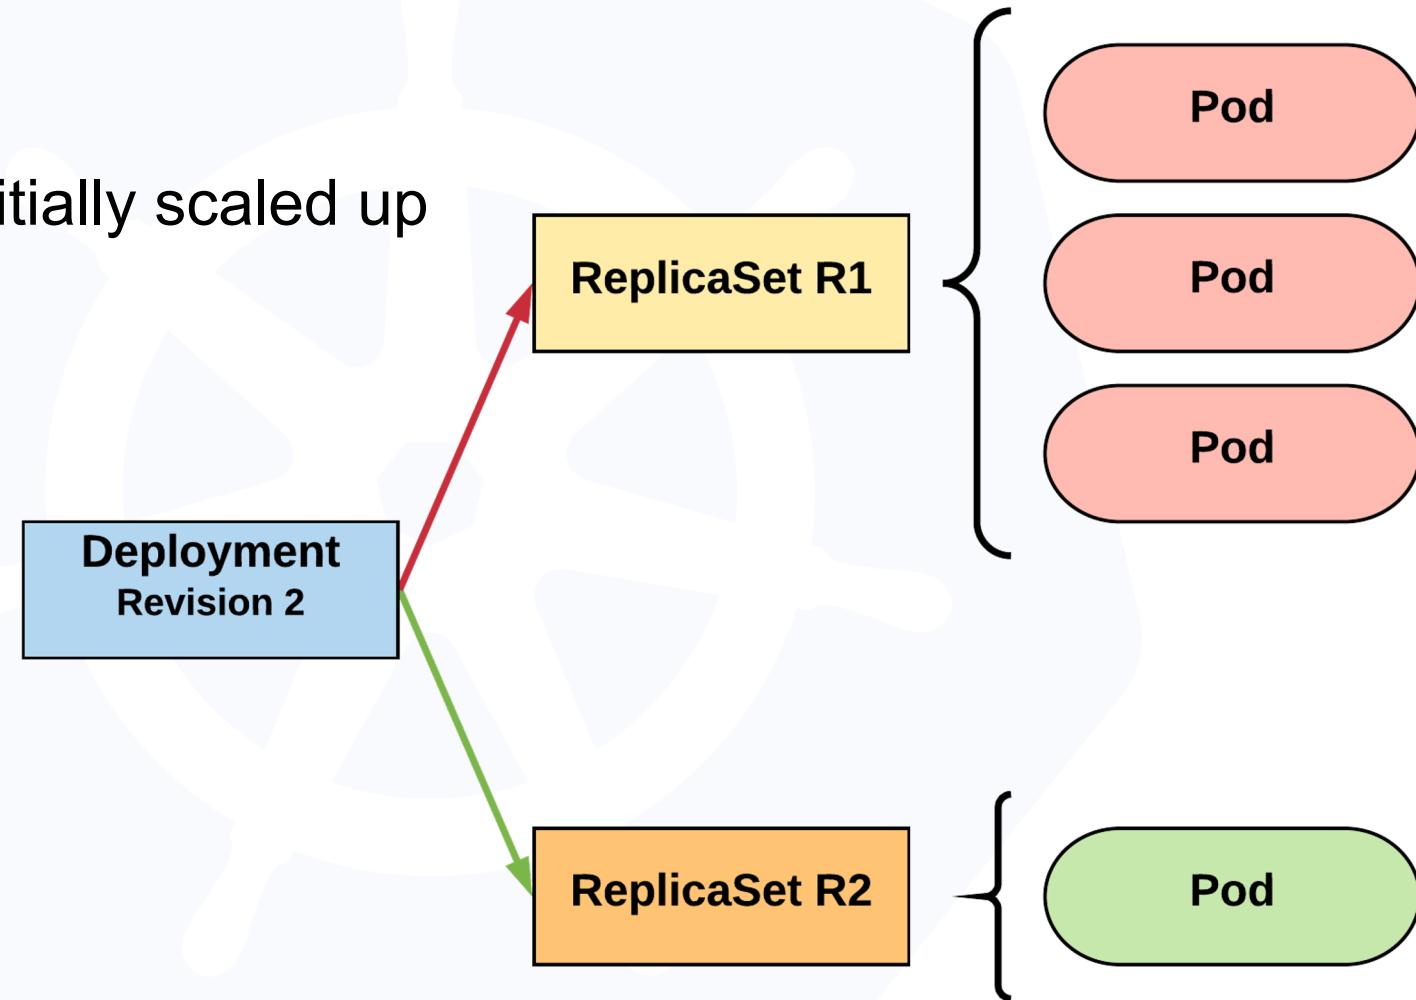




ROLLING UPDATE CONCEPT

New **ReplicaSet** is initially scaled up based on **maxSurge**.

R1 pod-template-hash:
676677fff
R2 pod-template-hash:
54f7ff7d6d

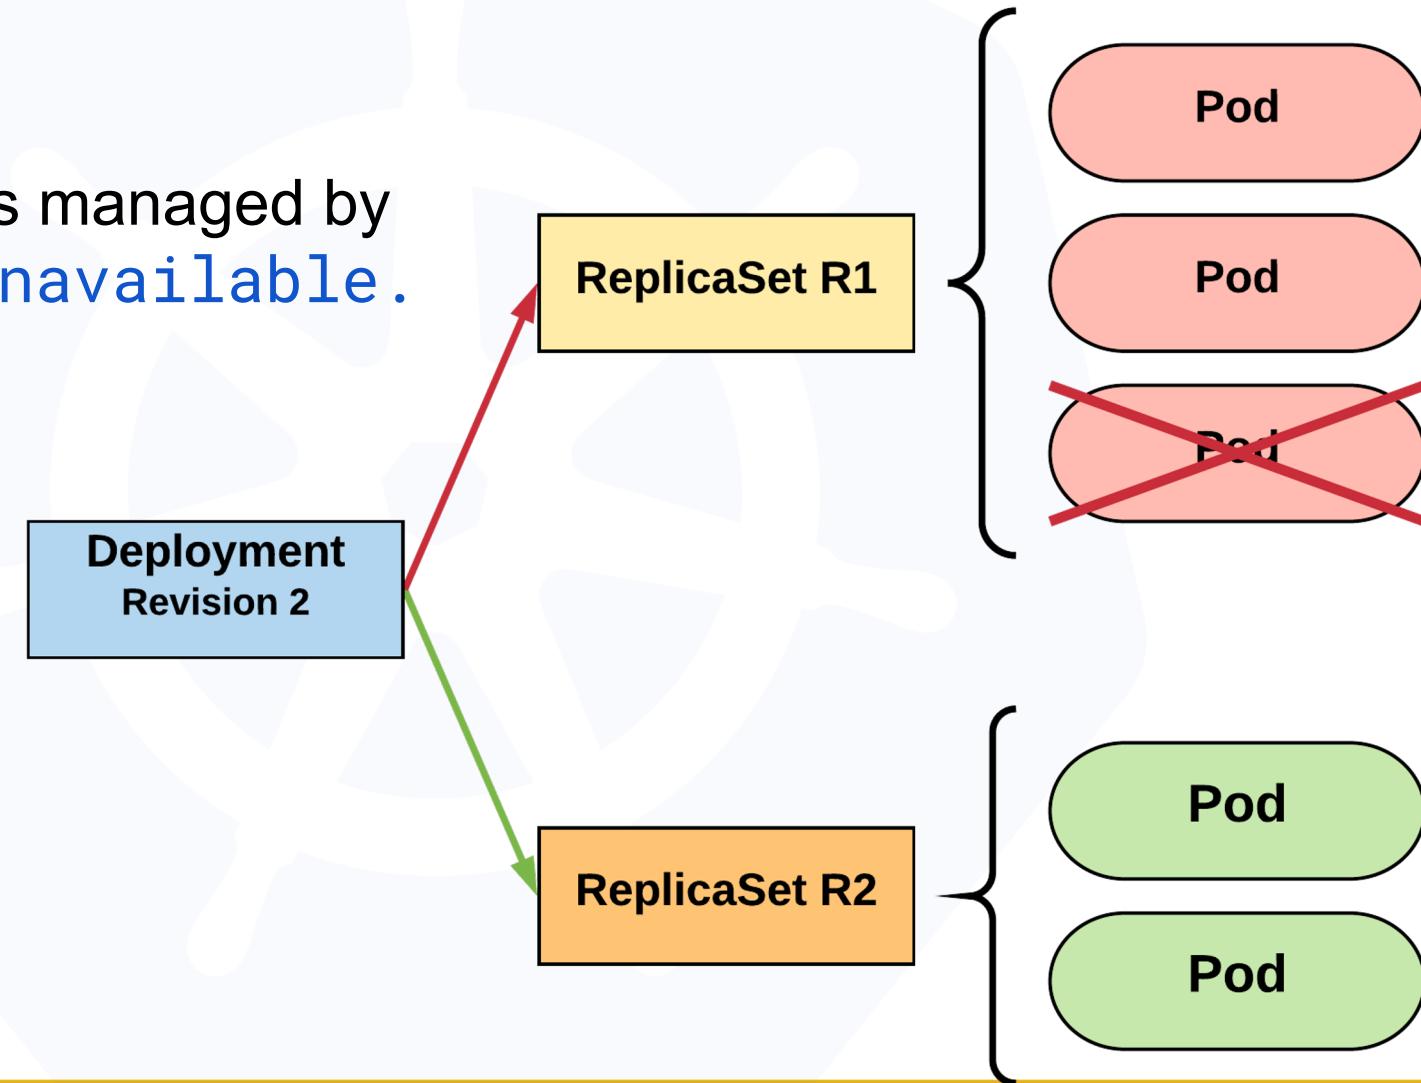




ROLLING UPDATE CONCEPT

Phase out of old Pods managed by **maxSurge** and **maxUnavailable**.

R1 pod-template-hash:
676677fff
R2 pod-template-hash:
54f7ff7d6d

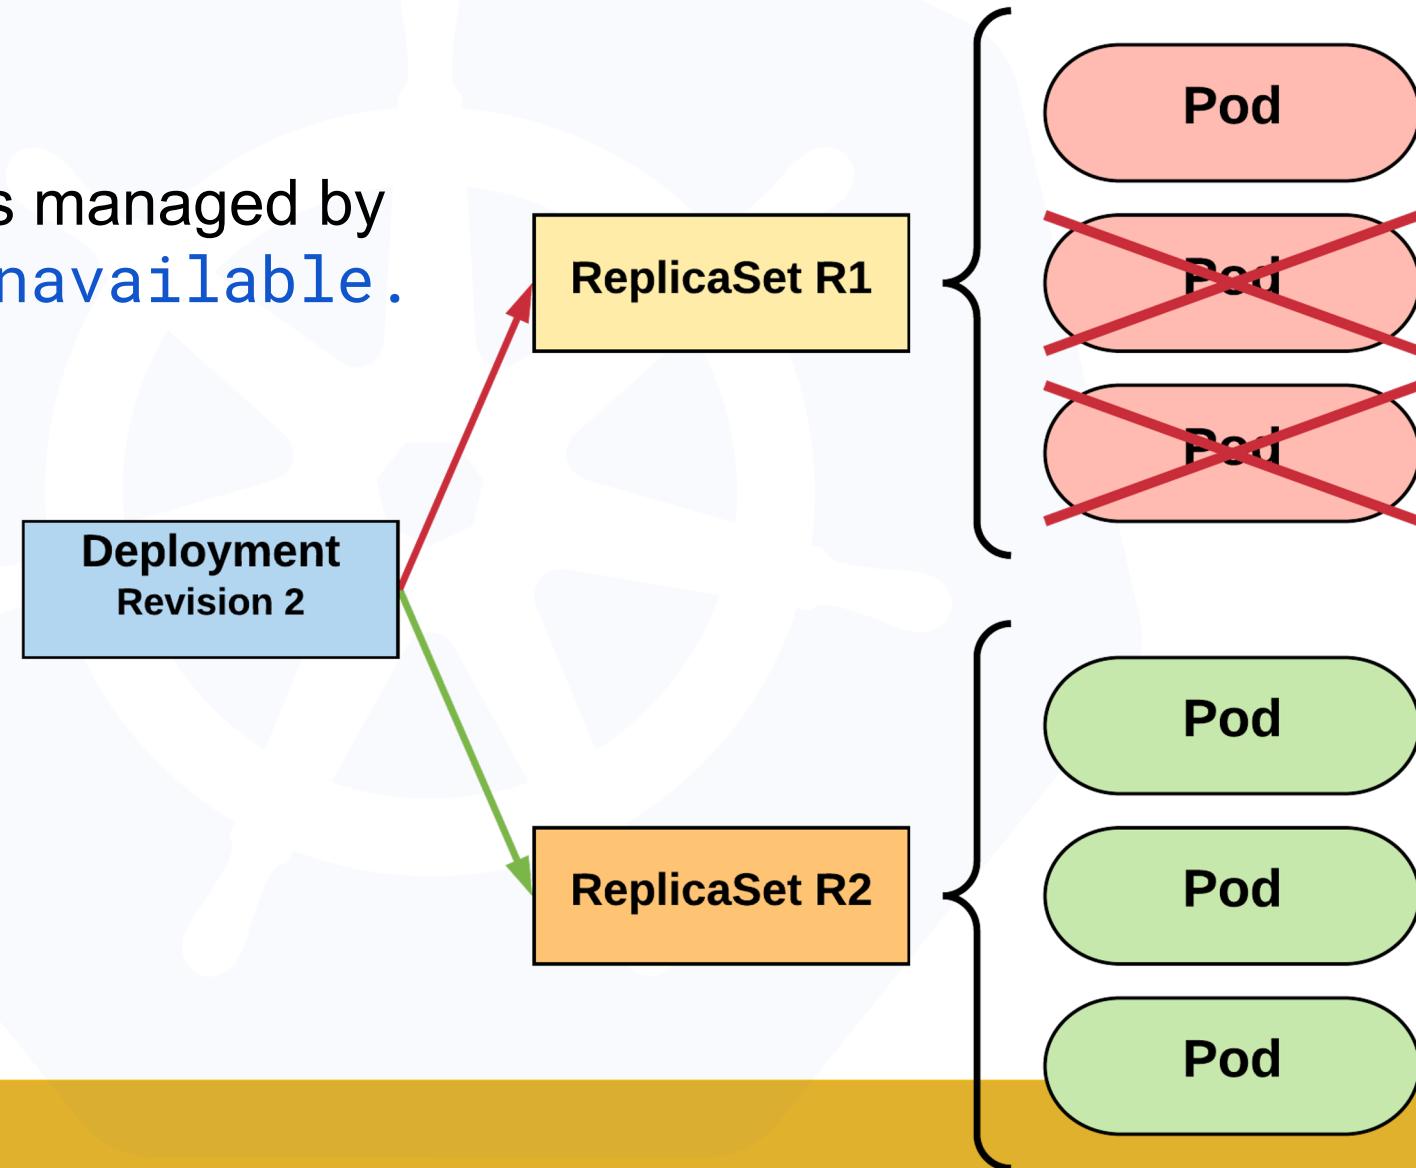




ROLLING UPDATE CONCEPT

Phase out of old Pods managed by **maxSurge** and **maxUnavailable**.

R1 pod-template-hash:
676677fff
R2 pod-template-hash:
54f7ff7d6d

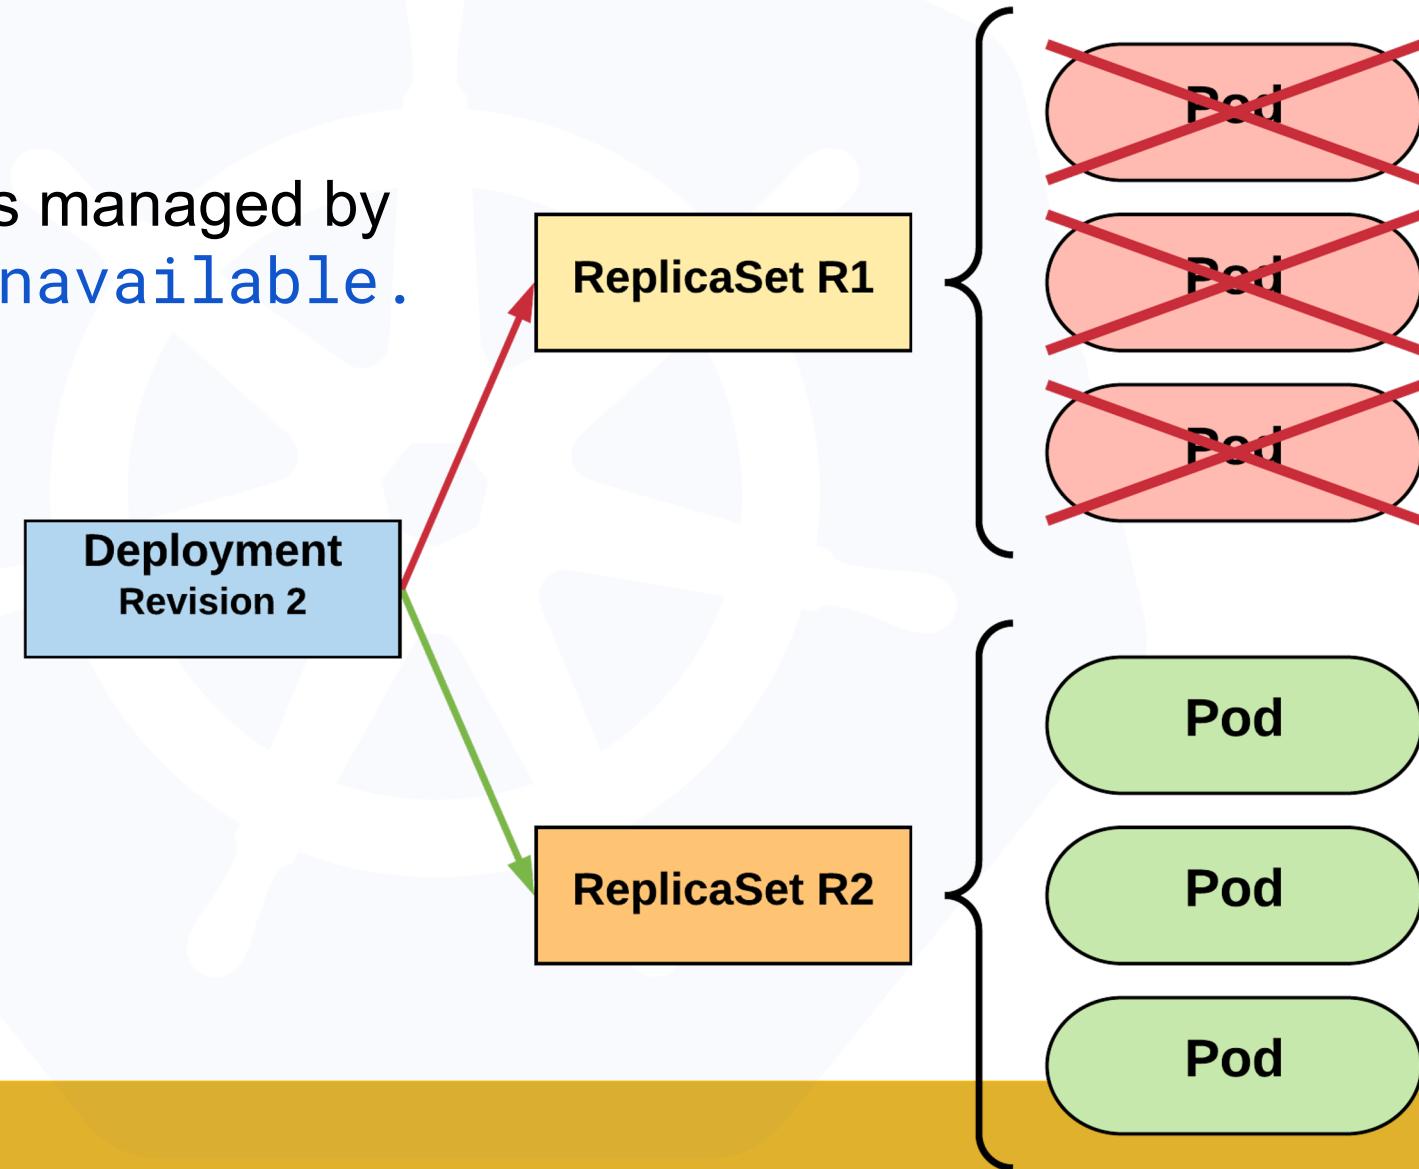




ROLLING UPDATE CONCEPT

Phase out of old Pods managed by **maxSurge** and **maxUnavailable**.

R1 pod-template-hash:
676677fff
R2 pod-template-hash:
54f7ff7d6d

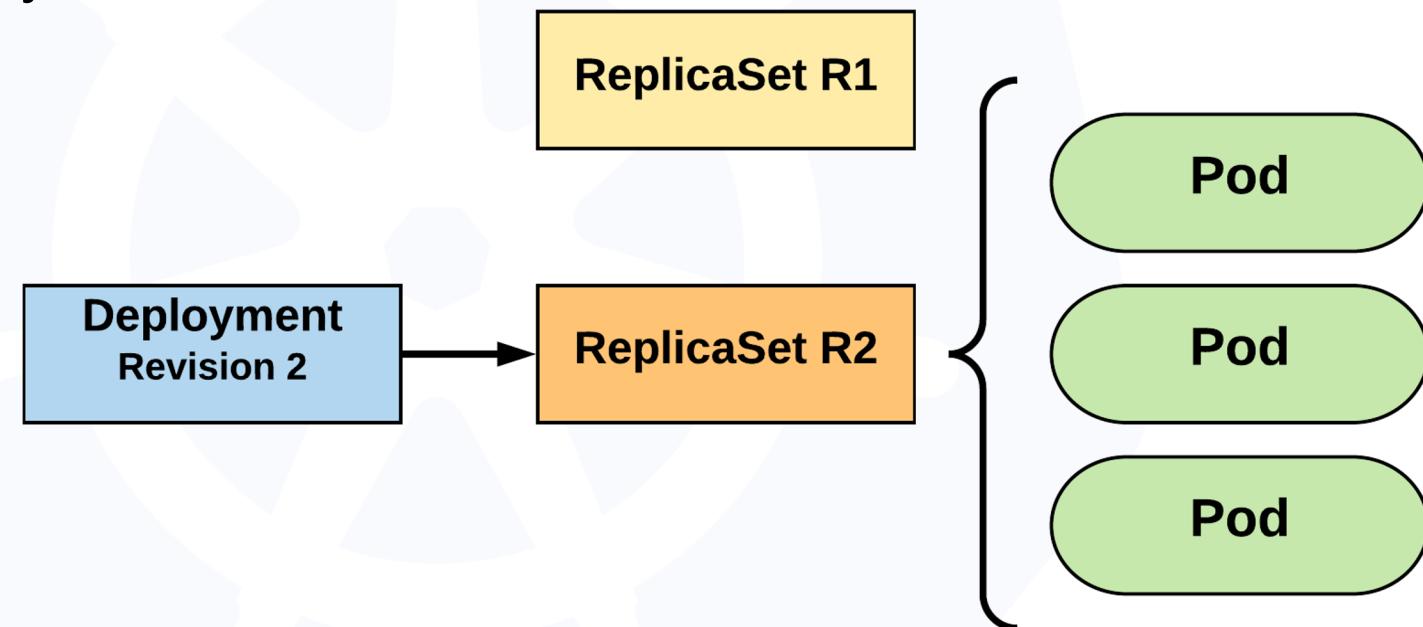




ROLLING UPDATE CONCEPT

Updated to new deployment revision completed.

R1 pod-template-hash:
676677fff
R2 pod-template-hash:
54f7ff7d6d





CONFIGURATION AND SENSITIVE DATA

Kubernetes has an integrated pattern for decoupling configuration from application or container.

This pattern makes use of two Kubernetes components: **ConfigMaps** and **Secrets**.

There are instances in your application where you want to be able to store runtime-specific information.

Things like connection info, SSL certs, or parameters that will probably change over time.



CONFIGURATION AND SENSITIVE DATA

CONFIGMAP: Configmap data is stored “outside” of the worker node (in etcd)

Externalized data stored within kubernetes.

Can be referenced through several different means:

- environment variable
- a command line argument (via env var)
- injected as a file into a volume mount

Can be created from a manifest, literals, directories, or files directly.



CONFIGURATION AND SENSITIVE DATA

All produce a **ConfigMap** with the same content

YAML

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: manifest-
example
data:
  city: Ann Arbor
  state: Michigan
```

CLI – from user input

```
$ kubectl create configmap literal-example \
> --from-literal="city=Ann Arbor" --from-literal=state=Michigan
configmap "literal-example" created
```

CLI – from file

```
$ cat info/city
Ann Arbor
$ cat info/state
Michigan
$ kubectl create configmap file-example --from-file=cm/city --
from-file=cm/state
configmap "file-example" created
```



CONFIGURATION AND SENSITIVE DATA

SECRET: Functionally identical to a ConfigMap.

Stored as base64 encoded content.

Encrypted at rest within etcd (if configured).

Ideal for username/passwords, certificates or other sensitive information that should not be stored in a container.

Can be created from a manifest, literals, directories, or from files directly.



CONFIGURATION AND SENSITIVE DATA

All produce a **Secret** with the same content

YAML (need to insert in base64)

```
apiVersion: v1
kind: Secret
metadata:
  name: manifest-example
type: Opaque
data:
  username: ZXhhbXBsZQ==
  password:
bXlwYXNzd29yZA==
```

CLI – from user input

```
$ kubectl create secret generic literal-secret \
> --from-literal=username=example \
> --from-literal=password=mypassword
secret "literal-secret" created
```

CLI – from file

```
$ cat secret/username
example
$ cat secret/password
mypassword
$ kubectl create secret generic file-secret --from-file=secret/username --
from-file=secret/password
Secret "file-secret" created
```



CONFIGURATION AND SENSITIVE DATA

Example of how to invoke

```
apiVersion: batch/v1
kind: Job
metadata:
  name: cm-env-example
spec:
  template:
    spec:
      containers:
        - name: mypod
          image: alpine:latest
          command: ["/bin/sh", "-c"]
          args: ["printenv CITY"]
      env:
        - name: CITY
          valueFrom:
            configMapKeyRef:
              name: manifest-example
              key: city
  restartPolicy: Never
```

```
apiVersion: batch/v1
kind: Job
metadata:
  name: secret-env-example
spec:
  template:
    spec:
      containers:
        - name: mypod
          image: alpine:latest
          command: ["/bin/sh", "-c"]
          args: ["printenv USERNAME"]
      env:
        - name: USERNAME
          valueFrom:
            secretKeyRef:
              name: manifest-example
              key: username
  restartPolicy: Never
```



RESOURCES

Kubernetes tutorial

<https://kubernetes.io/docs/tutorials/kubernetes-basics/>

Introduction to container orchestration

<https://www.exoscale.ch/syslog/2016/07/26/container-orch/>

TNS Research: The Present State of Container Orchestration

<https://thenewstack.io/tns-research-present-state-container-orchestration/>

Large-scale cluster management at Google with Borg

<https://research.google.com/pubs/pub43438.html>

Free Kubernetes Courses

<https://www.edx.org/>

Interactive Kubernetes Tutorials

<https://www.katacoda.com/courses/kubernetes>

Learn Kubernetes the Hard Way

<https://github.com/kelseyhightower/kubernetes-the-hard-way>

Official Kubernetes Youtube Channel

<https://www.youtube.com/c/KubernetesCommunity>

Official CNCF Youtube Channel

<https://www.youtube.com/c/cloudnativefdn>

Awesome Kubernetes

<https://ramitsurana.gitbooks.io/awesome-kubernetes/content/>



Q&A

Thanks
for your attention!
Questions?

Federico Accetta

federico_accetta@it.ibm.com