



UNIVERSITÀ  
CATTOLICA  
del Sacro Cuore

# NoSQL databases



# Definition

- “ A **NoSQL** (originally referring to "non-SQL" or "non-relational", now "Not Only" SQL) database provides a mechanism for storage and retrieval of data that is modeled in means other than the tabular relations used in relational databases.
- Motivations for this approach include: simplicity of design, simpler "horizontal" scaling to clusters of machines (which is a problem for relational databases), finer control over availability and limiting the object-relational impedance mismatch.

source: [NoSQL – Wikipedia](#)



# Traditional RDBMS

In the computing system (web and business applications), there are enormous data that comes out every day from the web. A large section of these data is handled by Relational database management systems (RDBMS). The idea of relational model came with E.F.Codd's 1970 paper "*A relational model of data for large shared data banks*" which made data modeling and application programming much easier. Beyond the intended benefits, the relational model is well-suited to client-server programming and today it is predominant technology for storing **structured data** in web and business applications.



# ACID rules

A database transaction, must be atomic, consistent, isolated and durable. Below we have discussed these four points.

- **Atomic** : A transaction is a logical unit of work which must be either completed with all of its data modifications, or none of them is performed.
- **Consistent** : At the end of the transaction, all data must be left in a consistent state.
- **Isolated** : Modifications of data performed by a transaction must be independent of another transaction. Unless this happens, the outcome of a transaction may be erroneous.
- **Durable** : When the transaction is completed, effects of the modifications performed by the transaction must be permanent in the system.

Often these four properties of a transaction are acronymed as **ACID**.



# CAP theorem

In theoretical computer science, the **CAP theorem**, also named **Brewer's theorem** after computer scientist Eric Brewer, states that it is impossible for a distributed data store to simultaneously provide more than two out of the following three guarantees:

- *Consistency*: Every read receives the most recent write or an error
- *Availability*: Every request receives a (non-error) response, without the guarantee that it contains the most recent write
- *Partition tolerance*: The system continues to operate despite an arbitrary number of messages being dropped (or delayed) by the network between nodes

When a network partition failure happens should we decide to

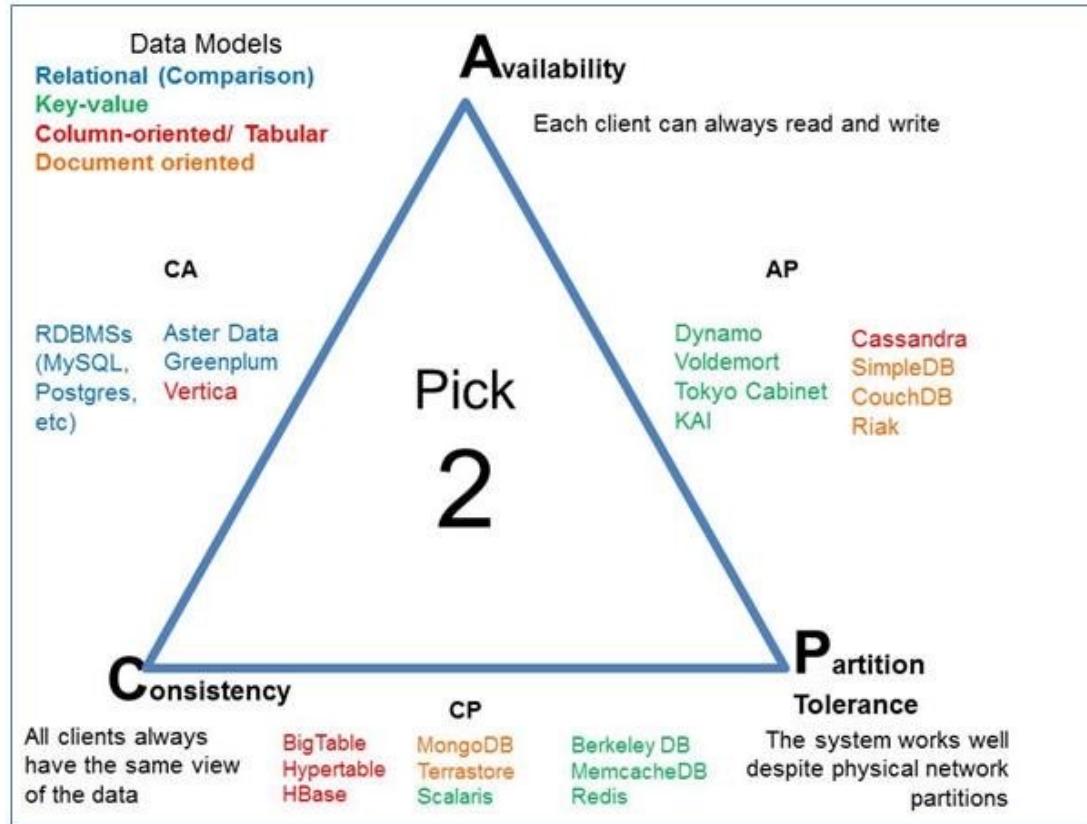
- Cancel the operation and thus decrease the availability but ensure consistency
- Proceed with the operation and thus provide availability but risk inconsistency

The CAP theorem implies that in the presence of a network partition, one **has to choose between consistency and availability**.

*Note that consistency as defined in the CAP theorem is quite different from the consistency guaranteed in ACID database transactions.*



# CAP theorem visualization





# BASE system

The CAP theorem states that a distributed computer system cannot guarantee all of the following three properties at the same time: Consistency, Availability, Partition tolerance.

A BASE system gives up on consistency.

- Basically Available indicates that the system does guarantee availability, in terms of the CAP theorem.
- Soft state indicates that the state of the system may change over time, even without input. This is because of the eventual consistency model.
- Eventual consistency indicates that the system will become consistent over time, given that the system doesn't receive input during that time.

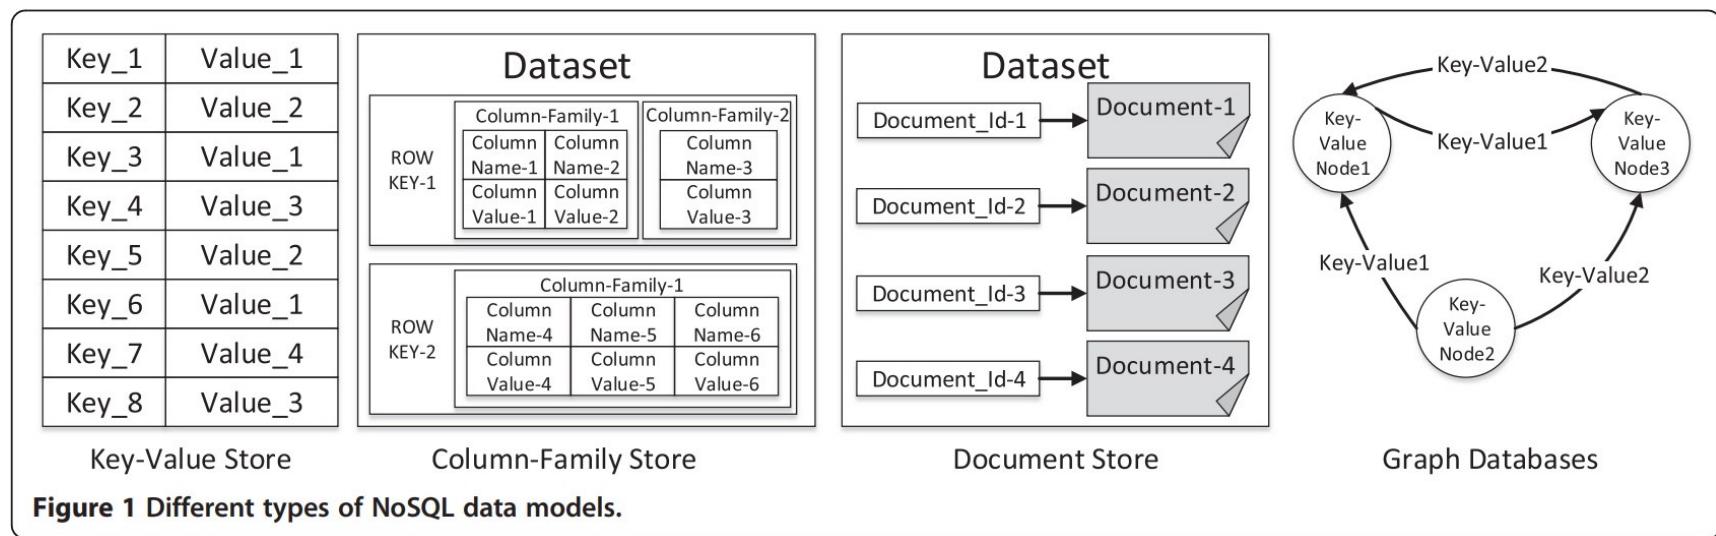


# NoSQL categories (data models)

There are four most common categories of NoSQL databases:

- Key-value stores
- Column-oriented
- Graph
- Document oriented

Each of these categories has its own specific attributes and limitations. There is not a single solution which is better than all the others, however there are some databases that are better to solve specific problems.



source: Grolinger et al., [Data management in cloud environments: NoSQL and NewSQL data stores](#)



# Key-value stores

- Key-value stores are most basic types of NoSQL databases.
- Designed to handle huge amounts of data.
- Based on Amazon's Dynamo paper.
- Key value stores allow developer to store schema-less data.
- In the key-value storage, database stores data as hash table where each key is unique and the value can be string, JSON, BLOB (Binary Large OBjec) etc.
- A key may be strings, hashes, lists, sets, sorted sets and values are stored against these keys.
- For example a key-value pair might consist of a key like "Name" that is associated with a value like "Robin".
- Key-Value stores can be used as collections, dictionaries, associative arrays etc.
- Key-Value stores follow the 'Availability' and 'Partition' aspects of CAP theorem.
- Key-Values stores would work well for shopping cart contents, or individual values like color schemes, a landing page URI, or a default account number.

**Example of Key-value store DataBase :** Redis, Dynamo, Riak. etc.



# Column-oriented

- Column-oriented databases primarily work on columns and every column is treated individually.
- Values of a single column are stored contiguously.
- Column stores data in column specific files.
- In Column stores, query processors work on columns too.
- All data within each column datafile have the same type which makes it ideal for compression.
- Column stores can improve the performance of queries as it can access specific column data.
- High performance on aggregation queries (e.g. COUNT, SUM, AVG, MIN, MAX).
- Works on data warehouses and business intelligence, customer relationship management (CRM), Library card catalogs etc.

**Example of Column-oriented databases :** BigTable, Cassandra, SimpleDB etc.



# Graph databases

- A graph database stores data in a graph.
- It is capable of elegantly representing any kind of data in a highly accessible way.
- A graph database is a collection of nodes and edges
- Each node represents an entity (such as a student or business) and each edge represents a connection or relationship between two nodes.
- Every node and edge are defined by a unique identifier.
- Each node knows its adjacent nodes.
- As the number of nodes increases, the cost of a local step (or hop) remains the same.
- Index for lookups.
- **Example of Graph databases :** OrientDB, Neo4J, Titan.etc.



# Document-oriented

- A collection of documents
- Data in this model is stored inside documents.
- A document is a key value collection where the key allows access to its value.
- Documents are not typically forced to have a schema and therefore are flexible and easy to change.
- Documents are stored into collections in order to group different kinds of data.
- Documents can contain many different key-value pairs, or key-array pairs, or even nested documents.

**Example of Document Oriented databases :** MongoDB, CouchDB etc.



# Performance

Data model	Performance	Scalability	Flexibility	Complexity	Functionality
Key-value store	high	high	high	none	variable (none)
Column-oriented store	high	high	moderate	low	minimal
Document-oriented store	high	variable (high)	high	low	variable (low)
Graph database	variable	variable	high	high	graph theory
Relational database	variable	variable	low	moderate	relational algebra

source: [NoSQL – Wikipedia](#)



# References

- [NoSQL introduction - w3resource](#)
- [What is NoSQL? NoSQL Databases Explained | MongoDB](#)
- [What is NoSQL? | Nonrelational Databases, Flexible Schema Data Models | AWS \(amazon.com\)](#)
- [CAP theorem – Wikipedia](#)



# Reading list

K. Grolinger, W.A. Higashino, A. Tiwari and M. Capretz, Data management in cloud environments: NoSQL and NewSQL data stores, *Journal of Cloud Computing*, 2013.

(documents available on Blackboard)



# Questions, comments?